

PREBAS: Preemptive Bandwidth Scaling for WebRTC in LEO Satellite Networks Using Lightweight Neural Prediction

Sushant Poudel

Dept. of Computer Science and Engineering
Nepal Engineering College, Bhaktapur, Nepal
sushant.poudel2028@gmail.com

Abstract—Low Earth Orbit (LEO) mega-constellations promise universal broadband, but their orbital mechanics — terminals moving at 7.6 km/s, with handoffs recurring every few minutes — produce abrupt capacity crashes that violate the stationary-channel assumptions behind modern real-time protocols. Google Congestion Control (GCC), the standard for Web Real-Time Communication (WebRTC), only detects congestion after a queue has already built, so during a LEO handoff the encoder keeps pushing a high bitrate into a channel that has already collapsed, producing severe bufferbloat and, in a typical five-minute session, more than 1,000 frame freezes. This paper introduces PREBAS (PREemptive BANDwidth Scaling), an edge-native predictive framework built around an ultra-lightweight one-dimensional convolutional neural network (1D-CNN). Reading a 320 ms sliding window of capacity, round-trip time (RTT), and loss-burst telemetry, the model isolates the causal “RTT creep” signature that appears 200–500 ms before a handoff, then preemptively scales bitrate down by a factor of $S_f = 0.40$ to drain the transmission queue before the crash, with a 300 ms cooldown preventing GCC from immediately reversing the cut. Across 30 independent stochastic trials — covering handoffs, ITU-R P.618-14 rain fade, and channel noise — PREBAS reduces frame freezes by 79% relative to GCC at a cost of only 8.9% of average bitrate, lifting the ITU-T P.910 QoE score by 8.36 points (43% relative) to within 2.9% of the Oracle ceiling. An ablation confirms the result depends on reading capacity, RTT, and loss jointly, filtering 92% of rain-fade false positives a single-signal heuristic misses. At 1,537 parameters and 0.091 ms inference latency on commodity ARM hardware, the model is deployable on existing edge infrastructure without GPU acceleration.

Index Terms—WebRTC, Low Earth Orbit (LEO) satellite networks, congestion control, predictive bandwidth scaling, edge artificial intelligence, convolutional neural network, bufferbloat mitigation, time-series classification, non-terrestrial networks (NTN), quality of experience (QoE).

I. Introduction

The convergence of Low Earth Orbit (LEO) satellite mega-constellations with terrestrial networks is redefining global real-time telecommunications. Deployments such as Starlink, OneWeb, and Kuiper are scaling to tens of thousands of satellites, increasingly tasked with delay-sensitive, high-bandwidth applications, notably Real-Time Communications (RTC) via the Web Real-Time Communication (WebRTC) protocol. As municipal AI kiosks, disaster-recovery networks, telemedicine platforms, and universal access initiatives expand into topologically challenging environments, continuous high-fidelity video over LEO links has become a genuine operational requirement, one that runs into a physical problem LEO cannot engineer around: unlike geostationary (GEO) systems, LEO satellites orbit at 300–2,000 km with orbital periods of roughly 90 to 120 minutes, staying overhead only briefly before handing off, and each handoff brings a sudden reconfiguration (new beams, new routes, Doppler-driven fading) that can crater

capacity almost instantly. Real-time protocols assume a channel degrading gradually enough for feedback loops to keep up.

WebRTC's default answer, Google Congestion Control (GCC), is fundamentally reactive: it scales bitrate only after Transport-wide Congestion Control (TWCC) feedback or queue-delay measurements reveal congestion has already occurred. That works on a terrestrial link, where congestion builds gradually, but fails mid-handoff: the encoder keeps pushing a steady-state bitrate into an already-collapsed pipe, the buffer fills before GCC notices, and severe bufferbloat, packet loss, and multi-second freezes follow, long enough to break a telemedicine call or emergency-response link outright.

The obvious next move is prediction, and the literature has tried it: LSTM networks, attention-based Transformers, and other sequence-learning approaches applied to channel telemetry. These models can see degradation coming but were not designed for a satellite terminal's compute budget; many carry millions of

parameters and tens to hundreds of milliseconds of inference time, untenable when the actionable window is a few hundred milliseconds. Static RTT thresholds run cheaply but are too blunt, unable to distinguish a genuine handoff from rain fade or a routing hiccup. What is missing is an edge-native model small enough for the terminal itself yet discriminating enough to recognize a handoff before it hits.

This paper addresses that gap with PREBAS, Preemptive Bandwidth Scaling, built around a minimal one-dimensional convolutional neural network (1D-CNN) watching a short sliding window of RTT, jitter, inter-arrival time, and loss-burst statistics. A LEO handoff announces itself 200–500 ms in advance through a subtle but consistent pattern (RTT creep, tightening jitter) statistically distinguishable from ordinary noise. When PREBAS recognizes that pattern, it scales the target bitrate down early, draining the send queue before capacity disappears rather than after.

This paper makes four contributions. First, we show through statistical analysis of synthetic and trace-driven channel data that LEO handoffs leave a learnable pre-event signature (RTT elevation, jitter amplification, loss-burst structure) separable from stochastic noise. Second, we design a predictive engine under 1,600 trainable parameters with sub-millisecond inference on ordinary ARM Cortex-A72 hardware, no GPU required. Third, we build and release a simulation environment modeling queue dynamics, handoff timing, and atmospheric fading for reproducible comparison between reactive and predictive control. Fourth, across the rain-fade, channel-noise, and Sf sensitivity scenarios in Section IV, we show preemptive queue clearance cuts WebRTC frame freezes by 79% and lifts ITU-T P.910 utility scores by 8.36 points (43% relative) over standard GCC, with minimal false-positive suppression on stable channels. Generalization across orbital inclinations and mobility profiles beyond the single geometry modeled here is future work (Section V-D).

Section II reviews related work; Section III presents the methodology; Section IV covers results, ablations, and sensitivity analysis; Section V discusses implications and limitations; Section VI concludes.

II. Related Work

Getting real-time video to survive a LEO handoff sits at the crossroads of congestion control, satellite measurement, and ML-based network prediction. What is still missing is a treatment of WebRTC's specific vulnerability to periodic, largely deterministic handoff events.

A. Reactive Congestion Control in RTC

Nearly every production RTC stack leans on reactive, feedback-driven control. GCC pairs a delay-based controller (Kalman-filtered inter-packet arrival times) with a loss-based controller reacting to RTCP loss reports [1], [4], [17]; SReAM [2] takes a similar posture. This works on terrestrial links, where congestion builds gradually, but poorly on a link whose capacity can fall off a cliff in under a second. Carlucci et al. [4] showed 95th-percentile queuing delay exceeding 195 ms under sudden capacity drops, and De Cicco et al. [15] found GCC needs roughly 25 seconds to converge after a capacity change. By the time GCC notices a handoff, the queues are already full.

B. Transport Challenges in LEO Constellations

Empirical measurement of Starlink, OneWeb, and Kuiper has grown rapidly. Zhao et al. [5] found stable OneWeb latency through most handovers but large swings at portal handoffs. Farooq et al. [6] found a roughly 15-second Starlink handover cadence with 140/75 ms boundary spikes; a more recent campaign [21] pins that schedule to specific seconds per minute, evidence the signature is structured rather than merely periodic. 3GPP Release 17/18 [7] targets the physical-layer handover but leaves queue buildup untouched: Valentine and Parisi [8] found RF handoffs push TCP Cubic's RTT from ~20 ms to over 61 ms, and Zhu et al. [16] showed neither loss- nor delay-based TCP reliably separates a path change from congestion. Physical-layer fixes alone cannot prevent this.

C. Predictive ML for Video Streaming, Edge Routing, and Satellite RTC

Pensieve [9], a DRL/A3C agent selecting DASH bitrates, reports 12–25% QoE gains, but DASH's ~4 s chunk buffering gives far more slack than WebRTC's sub-200 ms budget [10]. Zhang et al. [11] applied a Bi-LSTM for 3G/4G/5G prediction, and transformer forecasting shows promise in O-RAN [12]; both run tens of thousands to millions of parameters, too heavy for satellite terminals [13]. Static thresholds trade accuracy for cost: Ka-band rain attenuation exceeds 50 dB for 99.99% availability in some deployments [14], enough to trip a fixed threshold routinely.

Most relevant is Gottipati and Qiu [18], training a transformer policy to resize WebRTC's pacing queue from orbital ephemeris (SGP4/TLE) over two-minute windows, reporting up to 3× bitrate improvement and 62% fewer freezes in emulated sparse constellations; [18] remains a preprint under review, and we compare against its reported results. PREBAS instead reasons entirely from local telemetry (RTT, jitter, inter-arrival

time, loss bursts) over a much shorter 200–500 ms window, needing no ephemeris data or GPU. Zhang et al.'s multipath QUIC scheduler [20] and Zhao and Pan's LENS dataset [19] round out this space, which has not converged on a WebRTC-specific edge-native solution.

D. The Existing Gap

Reactive protocols induce bufferbloat by detecting congestion only after queues fill [1], [4]. Predictive ML for VoD or generic forecasting is usually too heavy for edge deployment [11]–[13]. Static thresholds cannot reliably separate a handoff from rain fade [14]. The newest handover-aware work reasons over minute-scale windows with external orbital data [18] or targets queue sizing and QUIC scheduling [18], [20], rather than WebRTC's millisecond-scale bitrate decision. That is the gap this paper addresses: an ultra-lightweight 1D-CNN on local telemetry alone, catching a handoff's signature early enough to act without extra compute or external data.

III. System Model and Methodology

Testing a claim like “predictive scaling beats reactive congestion control on LEO links” requires more than a toy simulation: a channel model that behaves like a real satellite link (noisy, occasionally deceptive, not conveniently periodic), plus a transport queue that responds the way a real WebRTC sender would. This section builds a stateful, discrete-time simulation pairing a stochastic NTN channel emulator with a physics-grounded transport queue, so neural preemption and reactive control can be compared on equal footing.

A. Stochastic NTN Channel Emulation

A GEO link barely moves; a LEO link is never still, its capacity swinging with orbital geometry, elevation-dependent path loss, and atmospheric conditions. Modeling handoffs as a fixed, repeating sequence would let the 1D-CNN memorize the pattern rather than learn it. The channel generator is built around three regimes grounded in satellite-channel theory.

1) *Steady State*: Baseline capacity $C(t)$ drifts under Ornstein–Uhlenbeck (mean-reverting) Gaussian noise:

$$C(t) = C_{\text{base}} + \varepsilon(t), \quad \varepsilon(t) \sim N(0, \sigma_c^2) \cdot e^{-\theta \Delta t}$$

with $C_{\text{base}} = 50$ Mbps, $\sigma_c = 2.5$ Mbps, mean-reversion rate $\theta = 0.1$, packet loss $p_{\text{loss}} = 0.005$, and base RTT $R_{\text{base}} = 35$ ms. This captures the slow, self-correcting drift of a stable LEO link, avoiding the flat white-noise assumption simpler generators use.

2) *Handoff Event and Causal Signature*: Handoff intervals are drawn from a Gamma distribution $\Gamma(k=20, \theta=25)$, giving a mean inter-handoff period of 500 steps (≈ 10 s at 50 Hz) and a coefficient of variation of 0.22, a

better fit for a positive-support renewal process than a Gaussian.

As elevation angle ε drops below $\varepsilon_{\text{min}} = 20^\circ$, the link shows a probabilistic “RTT creep,” the signature this paper targets. Path loss follows the standard slant-range model:

$$L_{\text{path}}(\varepsilon) = L_{\text{fs}} + 10 \log_{10}(d(\varepsilon)/d_0)^\alpha + \zeta_{\text{shadow}}$$

where L_{fs} is free-space loss, $d(\varepsilon) = h_s/\sin(\varepsilon)$ is slant distance at altitude $h_s = 550$ km, $\alpha = 2.5$, and $\zeta_{\text{shadow}} \sim N(0, \sigma_s^2)$ models shadowed-Rician fading. As $\varepsilon \rightarrow \varepsilon_{\text{min}}$, RTT climbs by up to $\Delta R_{\text{max}} = +15$ ms from ingress-buffer buildup, the pre-event window PREBAS targets.

The make-before-break switch then drops capacity to a floor of $C_{\text{floor}} \approx 1.5$ Mbps, with routing delays exceeding $R_{\text{handoff}} > 80$ ms as BGP reconverges.

3) *Rain Fade (False Positive Injection)*: Precision is stress-tested against attenuation events from the ITU-R P.618-14 rain fade model: capacity drops up to 20 dB at Ka-band, but without RTT degradation. Rain fade is a two-state Markov chain (clear/rain) tuned for tropical climates ($P_{\text{rain}} = 0.15$, mean duration = 120 s), forcing the classifier to separate the correlated RTT–jitter–loss triplet of a true handoff from capacity-only interference.

The channel state at time t is the tuple $S(t) = \{C(t), R(t), L(t), \phi(t)\}$, $\phi(t) \in \{\text{steady, handoff, rain}\}$.

B. Edge-Native Neural Predictor (1D-CNN)

Running deep learning on a municipal edge router or terminal user equipment involves hard constraints that most “lightweight” models still exceed. LSTMs are the conventional time-series choice, but their parameter count and sequential graph are a poor match for a 20 ms WebRTC feedback loop. Recent benchmarking confirms 1D-CNNs reach comparable accuracy while needing far less RAM and flash: 27.6 ms latency versus upward of 2 seconds for an equivalent LSTM under INT8 quantization.

The predictor here is purpose-built for sub-millisecond inference rather than adapted from something heavier. It processes a sliding window $X_t \in \mathbb{R}^{W \times F}$ of $W = 16$ steps (320 ms at 50 Hz) across $F = 3$ features: estimated capacity $\hat{C}$, RTT R , and loss rate L . A grid search over $W \in \{8, 16, 32, 64\}$ showed $W = 16$ maximized validation F1-score while keeping latency below 0.1 ms; $W = 8$ under-captured the creep signature, and $W \geq 32$ added cost without a corresponding accuracy gain.

Architecture

The network uses standard 1-D convolutions with BatchNorm regularization for training stability. Table 1 lists the layer-by-layer structure and parameter accounting, verified by direct arithmetic.

TABLE 1. Network Architecture and Parameter Count

Layer	Type	Units/Filters	Kernel	Output Shape	Params
1	Conv1D	8	3	(14, 8)	80
2	BatchNorm1D	8	—	(14, 8)	16
3	MaxPool1D	—	2	(7, 8)	0
4	Conv1D	16	3	(5, 16)	400
5	BatchNorm1D	16	—	(5, 16)	32
6	MaxPool1D	—	2	(2, 16)	0
7	Flatten	—	—	32	0
8	Dense (ReLU)	23	—	23	759
9	Dense (ReLU)	10	—	10	240
10	Output (Sigmoid, no bias)	1	—	1	10

The total trainable parameter count is verified layer-by-layer as:

$$P_{\text{total}} = 80 + 16 + 400 + 32 + 759 + 240 + 10 = 1,537 \text{ parameters}$$

This 1,537-parameter count is used consistently throughout the paper (complexity analysis, results, conclusion) and is well below the 1,600-parameter budget, a greater than 100× reduction relative to the smallest LSTM that could plausibly handle this task. Both BatchNorm layers contribute only their trainable scale-and-shift parameters (2 per channel); running-mean/variance statistics are updated during training but not counted as trainable weights. No post-hoc calibration layer is used: the sigmoid output is trained end-to-end, and its decision threshold $\tau = 0.5$ (Section III-C) is verified as the operating point matching the precision-recall tradeoff reported in Section IV-A.

Benchmarking on an ARM Cortex-A72 (representative of edge-router SoCs) at 1.8 GHz shows per-step inference latency under 0.1 ms with TensorFlow Lite INT8 quantization, under 0.5% of the 20 ms WebRTC transmission budget. Memory footprint is 6.2 KB (INT8-quantized weights) plus 3.1 KB (activations), 9.3 KB total; this figure is used consistently in Sections IV-E, V-C, and VI.

Training Regime

The model trains on 50,000 stochastically generated channel traces (2,000 steps each) under binary cross-entropy loss, batch size 256, mixed-precision (FP16). A window is labeled positive if a handoff occurs within $T_{\text{pred}} = 6$ steps (120 ms) of its final timestep. Class imbalance is addressed via focal loss ($\gamma = 2.0$) and a 3:1 negative-to-positive sampling ratio, with Adam ($\eta = 0.001$, $\beta_1 = 0.9$, $\beta_2 = 0.999$) and early stopping (patience 3 epochs, 20% validation holdout). Best validation loss occurs at epoch 9; training runs 12 epochs total, roughly 4 hours on a single NVIDIA RTX 4090. All results use the epoch-9 checkpoint.

C. Preemptive Congestion Control Engine

The predicted signal only matters if it changes transport behavior. Standard GCC applies AIMD once its delay-based controller's Kalman-filtered gradient flags overuse. Target bitrate $A(t)$ updates as:

$$A(t) = \eta \cdot A(t-1) \text{ if } \sigma = \text{Increase}$$

$$A(t) = \alpha \cdot R_r(t) \text{ if } \sigma = \text{Decrease}$$

$$A(t) = A(t-1) \text{ if } \sigma = \text{Hold}$$

where $\eta = 1.05$, $\alpha = 0.85$, and $R_r(t)$ is the receiving rate. This loop needs roughly 100 ms to detect overuse and another 200–500 ms to drain a saturated queue, nowhere near fast enough for a LEO handoff.

When the classifier's handoff probability crosses $P_{\text{handoff}} \geq \tau = 0.5$, a preemptive suppression trigger fires immediately:

$$A_{\text{preempt}}(t) = S_f \cdot A(t), \quad S_f = 0.40$$

This 60% cut manufactures an artificial data deficit, draining the transmission queue $Q(t)$ per the fluid model:

$$dQ/dt = \lambda_{\text{in}}(t) - \mu_{\text{out}}(t)$$

where $\lambda_{\text{in}} = A_{\text{preempt}}$ and $\mu_{\text{out}} = C(t)$. When $C(t) \gg A_{\text{preempt}}$, the queue drains exponentially; when $C(t) \approx A_{\text{preempt}}$, it settles at a low watermark.

To stop GCC's reactive loop from immediately overriding this cut, the framework enforces a $T_{\text{cooldown}} = 15$ -step (300 ms) suppressive cooldown: the delay-based detector is gated to Hold, the loss-based controller's decrease factor is bounded to $\alpha_{\text{min}} = 0.60$, and AIMD recovery stays suspended until T_{cooldown} expires or the classifier drops below $P_{\text{handoff}} < 0.2$ (early release). This gives the preemptive action time to drain the queue while avoiding an indefinite lockout if the predicted handoff does not materialize.

Production Integration: The Neural Shim

The description above assumes direct override of GCC's internal target bitrate, appropriate for simulation but incomplete for production, since REMB and TWCC are distinct mechanisms owned by different endpoints. In the deployment we envision, the 1D-CNN runs at the sender alongside its own bandwidth estimator and, on a positive detection, adjusts the sender's target bitrate through the same internal API GCC uses for AIMD (e.g., `SetTargetBitrate()` in a Chromium-derived stack), requiring no receiver or wire-protocol change. It can additionally synthesize a REMB message reflecting the reduced target, an established pattern for keeping bidirectional estimators aligned. We call this pattern, predictive sender-side adjustment with optional REMB-consistent signaling, a “neural shim.” Full integration into a production codebase, including interception of receiver-generated TWCC feedback, is future work (Section V-D).

D. Benchmarking and Evaluation Rigor

The framework evaluates the architecture across $N = 30$ independent trials, each spanning 10,000 steps (200 s simulated time) with randomized seeds, so no two trials share a handoff sequence or rain-fade pattern. The controller is compared against three baselines: Reactive GCC (unmodified, default parameters, the industry-standard lower bound); an EWMA Heuristic ($\alpha = 0.25$ on RTT, triggering when smoothed RTT exceeds $R_{\text{base}} + 10$ ms); and an Oracle upper bound (perfect future knowledge with optimal S_f and T_{cooldown} , the ceiling any preemptive architecture could reach).

Performance uses a strict, ITU-T P.910-inspired logarithmic QoE utility function penalizing structural degradation over minor bandwidth reductions:

$$U_{\text{QoE}} = \alpha \log_{10}(B/B_{\text{max}}) - \beta L - \gamma(D/D_{\text{max}}) - \delta \cdot 1\{F\}$$

where B is average bitrate in kbps ($B_{\text{max}} = 50,000$ kbps), L is packet-loss fraction, D is 95th-percentile queuing delay in ms ($D_{\text{max}} = 500$ ms), and F indicates catastrophic freezes, defined as $D(t) > 175$ ms for at least 3 consecutive frames (≈ 100 ms at 30 fps). Weights $\alpha = 10.0$, $\beta = 5.0$, $\gamma = 2.0$, $\delta = 15.0$ reflect the relative severity of each failure mode. The logarithmic bitrate term reflects diminishing marginal utility of bandwidth (Weber–Fechner scaling); the freeze penalty δ is deliberately heavy since a conversational freeze does not degrade gracefully the way a bitrate dip does.

All metrics are mean $\pm$ standard error across $N = 30$ trials. Pairwise comparisons use the Wilcoxon signed-rank test with Holm–Bonferroni correction ($m = 3$); effect sizes use Cliff’s delta (δ_{cliff}). Significance is declared at $p < 0.01$.

IV. Results and Evaluation

With the simulation framework from Section III in place, the central question is whether this actually works end to end. We ran an empirical evaluation across $N = 30$ independent trials, each spanning 15,000 transmission steps (five minutes of simulated real-time flight at 50 Hz), with sufficient variation in handoff timing, rain fade, and background noise that no single trial resembles another. Unless noted, metrics are mean $\pm$ standard error; pairwise comparisons use the Wilcoxon signed-rank test with Holm–Bonferroni correction ($m = 3$), effect sizes via Cliff’s delta (δ_{cliff}).

A. Predictive Model Performance and Classification

Before transport metrics, it is worth isolating the classifier’s own performance, since the requirement is demanding: catch structural capacity crashes with high confidence while staying quiet through rain fade, congestion spikes, and channel noise. False positives matter especially, since every spurious trigger costs a

300 ms AIMD cooldown that eats into steady-state throughput.




Fig. 1. Confusion matrix for the 1D-CNN classifier on a held-out operational trace of 19,984 steps, combining true orbital handoffs with injected rain fade.

Fig. 1 shows specificity of 99.84%, with only 31 of 18,932 clean samples misclassified as handoffs, a false positive rate of 0.16%, the single most important number here: it confirms the model will not throttle bitrate on a whim during ordinary fluctuation. Precision is 93.80%, so when the model fires, it is correct nearly 94% of the time.

Recall, at 44.58% (469 of 1,052 true handoffs caught), looks low in isolation but reflects a deliberate choice: prioritizing confident detections over exhaustive coverage, defensible since a missed handoff falls back to ordinary GCC behavior while a false positive actively degrades QoE. On a class-balanced set, the architecture reaches an F1-score of 0.842 (precision 0.815, recall 0.871), confirming the lower operational recall stems from class imbalance (5.3% positive prevalence), not model incapacity.

The rain-fade injection proved an important design decision. A naïve RTT threshold ($R_{\text{base}} + 10$ ms) fires on 88% of rain-fade events, essentially useless in tropical deployments. The 1D-CNN instead learns the relationship between a capacity drop and the presence of RTT creep, correctly filtering 92% of weather anomalies, evidence it has learned a causal signature rather than a simple threshold.

B. End-to-End Transport Performance

Table 2 summarizes mean performance across all 30 trials with 95% bootstrap confidence intervals. Unless noted, “Total Freezes” uses the strict ITU-T P.910 definition (queuing delay above 175 ms sustained for at least 3 consecutive frames), aggregated per five-minute session.

TABLE 2. End-to-End Transport Metrics (N = 30 Trials)

Control Algorithm	Total Freezes ($\mu \pm \text{CI}$)	Avg. Bitrate (Mbps)	ITU-T QoE Score
Reactive GCC	1042.3 $\pm$ 68.5	32.41	19.45 $\pm$ 1.8
EWMA Heuristic	485.1 $\pm$ 42.2	31.12	23.18 $\pm$ 1.5

Preemptive Neural (Ours)	218.4 $\pm$ 25.1	29.54	27.81 $\pm$ 1.2
Oracle (Upper Bound)	175.2 $\pm$ 18.0	29.15	28.65 $\pm$ 1.1

Unmodified GCC saturates its queues badly during handoffs, averaging over 1,000 freeze events per session. The preemptive neural controller cuts that by 79% relative to GCC ($p < 0.001$, $\delta\text{cliff} = 0.93$) and 55.0% relative to the EWMA heuristic ($p < 0.001$, $\delta\text{cliff} = 0.71$), at a modest cost of 8.9% average bitrate, landing the QoE score at 27.81, an 8.36-point (43%) improvement over GCC's 19.45, within 2.9% of the Oracle ceiling of 28.65.




Fig. 2. Empirical cumulative distribution function of effective queueing latency across all 30 trials, with the ITU-T 175 ms freeze threshold marked.

Fig. 2 shows tail behavior a table cannot. All four curves rise steeply to roughly 0.90 by 50 ms; under quiescent conditions every algorithm looks similar, since none of this matters until a handoff hits. Past the 175 ms threshold, reactive GCC shows roughly 8% violation probability, the preemptive neural controller near 6%, tracking the Oracle bound closely. Its occasional missed detections (from 44.6% recall) do not translate into catastrophic tail degradation; GCC's curve has the flattest tail slope, reflecting bufferbloat recovery playing out repeatedly.

C. Feature Ablation and Causal Signature Verification

Claiming the model learned a causal signature rather than memorized noise requires demonstrating it. We ran a controlled feature ablation at a fixed $S_f = 0.40$, reported in Table 3 as freeze reduction relative to reactive GCC for each input configuration.

TABLE 3. Feature Ablation Study ($S_f = 0.40$, $N = 30$)

Input Features	Freeze Reduction vs. GCC	FPR during Rain Fade	Interpretation
Capacity only	3.2%	12%	Ambiguous; cannot separate handoff from fade
Loss only	1.8%	8%	Lagging indicator; reactive by nature
RTT only	65.0%	34%	Captures creep, but triggers on non-handoff congestion

Capacity + Loss + RTT (Full)	79.0%	0.8%	Multi-modal causal signature; best precision-recall tradeoff
------------------------------	-------	------	--

Capacity or loss alone performs nearly identically to the reactive baseline, confirming neither univariate signal carries enough predictive weight on its own. RTT alone reaches 65% freeze reduction, confirming RTT creep is the dominant feature, but with a 34% false positive rate too noisy for production use alone. Only the full three-channel model separates localized bufferbloat (RTT rises, capacity stable) from genuine physical degradation (capacity drop, RTT creep, and loss-burst structure together), showing the model exploits the correlated triplet rather than any single signal.

D. Hyperparameter Sensitivity and Pareto Optimization

Real-world behavior depends heavily on how aggressively bitrate is suppressed ahead of a predicted handoff, governed by $S_f \in [0.20, 0.90]$. Fig. 3 sweeps this parameter and plots the resulting Pareto frontier of throughput against total freeze count.




Fig. 3. Pareto frontier of average system throughput versus aggregate video frame freezes across the S_f sensitivity sweep, using a finer-grained freeze definition (sustained excursions above 100 ms, ≥ 2 frames) than the ITU-T 175 ms/3-frame threshold used in Table 2, chosen to resolve differences between closely spaced S_f values that the stricter threshold cannot separate.

Three regimes emerge. Push S_f too high (≥ 0.70) and the bitrate cut is insufficient to clear the queue before the crash; performance regresses toward EWMA-heuristic territory. Push it too low (≤ 0.20) and freeze counts fall slightly further, but the encoder starves for data, and the logarithmic bitrate term penalizes that starvation more than the freeze improvement is worth. The Pareto optimum sits in a narrow band around $S_f = 0.40$ – 0.45 ; below it, marginal freeze reduction is not worth the bitrate cost, and above it, freeze counts climb again as queue drainage becomes incomplete. Using the strict 175 ms/3-frame threshold from Table 2, we separately confirmed $S_f = 0.40$ remains freeze-minimizing within this band; the finer threshold in Fig. 3 changes only

resolution, not the optimum's location. We selected $Sf = 0.40$ as the production default.




Fig. 4. Representative two-handoff time-series sequence (steps 0–350, 20 ms/step) showing physical channel capacity, GCC and Preemptive Neural bitrate response (top), and resulting queue latency (bottom), at the selected operating point $Sf = 0.40$.

Fig. 4 zooms into one representative two-handoff sequence at this operating point. Physical capacity (dotted) crashes twice, and the controllers respond very differently. Around step 75, the neural controller detects the pre-handoff RTT creep and cuts bitrate to $Sf \cdot A = 0.40 \times 45 \approx 18$ Mbps before capacity actually drops. GCC has no such warning, holding 45 Mbps until roughly step 85, by which point the queue is already overfilled. The neural approach takes a brief queue spike to about 240 ms around step 75, then recovers to its under-40 ms baseline within about 10 steps; GCC, at the second handoff near step 190, lets the queue build monotonically, crossing 175 ms and staying there for more than 30 steps. This is the clearest single illustration of the paper's argument: the queue drains before the crash, not after it.

E. Computational Overhead and Edge Viability

None of the above matters if the model cannot run where it needs to. The architecture in Section III-B contains 1,537 trainable parameters, over $100\times$ smaller than the smallest LSTM that could plausibly handle this task. Benchmarking on a single ARM Cortex-A72 core at 1.8 GHz, representative of municipal-kiosk and edge-router SoCs, using TensorFlow Lite INT8 quantization over 10,000 continuous iterations, gave a mean inference latency of 0.091 ± 0.003 ms, 99th-percentile latency of 0.12 ms, a weight footprint of 6.2 KB (INT8) and activation footprint of 3.1 KB.

Since RTCP TWCC feedback and video frame generation both operate on a 20 ms cadence, an inference delay under 0.1 ms consumes less than 0.5% of the available budget, and even the 99th-percentile latency stays two orders of magnitude below the frame deadline.

The 1D-CNN adds no meaningful computational penalty, enabling real-time execution on constrained edge hardware without GPU acceleration.

These results come from a single-flow, single-orbital-plane simulation with idealized MAC-layer handoff signaling. Multi-flow coordination, multi-orbital-plane generalization, and validation against proprietary constellation hardware remain open questions, discussed as limitations in Section V-D.

V. Discussion

The results in Section IV support this paper's central claim: predictive, application-layer queue clearance can meaningfully offset the transport problems LEO orbital mechanics forces onto real-time protocols. This section examines the trade-offs, the design rationale, and the framework's actual limits.

A. The QoE Trade-off: Throughput vs. Temporal Continuity

GCC and its reactive relatives chase maximum utilization, operating at the edge of queue saturation. Fig. 3's Pareto frontier shows our architecture doing the opposite: deliberately giving up unbuffered throughput for structural stability. At $Sf = 0.40$, the system accepts an 8.9% bitrate cut to buy a 79% freeze reduction, a favorable trade since the two failure modes are not perceived equally. A drop from 1080p to 720p reads as a gradual, forgivable dip; a freeze past 175 ms breaks conversational rhythm outright. The ITU-T P.910 function encodes this directly: freeze penalty $\delta = 15.0$ sits well above bitrate term $\alpha = 10.0$. Mechanically, the 1D-CNN hides the handoff from perception: the encoder dips resolution briefly, but the stream never stalls.

B. Why Neural? The Necessity of Multi-Variate Causal Inference

A fair critique of ML-in-networking work is that a complex model often solves what a heuristic could handle. The ablation in Section IV-C answers that directly. A static RTT threshold reaches 65% freeze reduction but fires on 34% of ordinary congestion, an unacceptable rate. A capacity-slope heuristic fails more cleanly: rain fade produces a capacity drop indistinguishable, on that axis alone, from a real handoff. The 1D-CNN succeeds by reading joint structure across the window; a genuine handoff shows RTT creep, capacity slope, and loss-burst structure moving together, while rain fade shows isolated capacity loss with no delay shift, reflecting the underlying physics of elevation angle, slant distance, and propagation delay, a dependency chain no static threshold can encode.

C. Edge Viability: Breaking the Inference-Latency Barrier

Predictive transport control has historically stalled on inference latency: Pensieve needs well over 100 ms per policy evaluation [9], and transformer forecasting runs into millions of parameters with GPU dependence [12], both past WebRTC's sub-20 ms budget. The 1D-CNN sidesteps this: at 1,537 parameters and 0.091 ms latency on an ARM Cortex-A72, it consumes roughly 0.46% of a video frame interval, and its 6.2 KB INT8-quantized weights fit entirely in L1 cache. Running prediction on the terminal or first-hop router avoids the cloud round-trip that would reintroduce this latency; no FPGA, no GPU, deployable on hardware that already exists.

D. Limitations and Pathways to Deployment

The channel emulator models the LEO characteristics that matter — Doppler shift, routing-delay recalculation, fading calibrated against 3GPP TR 38.811 [7] — but remains a simulation, with three gaps worth stating directly.

1) *Proprietary PHY/MAC abstraction*: The emulator assumes idealized MAC-layer handoff signaling rather than modeling proprietary beam-management and HARQ procedures specific to Starlink, Kuiper, or OneWeb hardware. Validating against open-interface NTN testbeds, and integrating the neural shim into a production stack for measurement against live traffic, is the natural next step.

2) *Single-flow isolation*: The evaluation assumes one WebRTC flow per terminal. Multi-flow deployments would need Sf to respond to aggregate queue occupancy; a multi-agent extension with per-flow 1D-CNNs coordinating a shared Sf is plausible, though untested here.

3) *Integration with standardized feedback*: The neural shim (Section III-C) is designed to sit in front of unmodified WebRTC endpoints without client-side changes. Empirical validation inside a real browser stack is the immediate next phase of this work.

VI. Conclusion

LEO mega-constellations are one of the more credible paths toward closing the connectivity gap for the roughly half of humanity still without reliable broadband, but frequent handoffs producing capacity crashes reactive congestion control cannot see coming cost more than they might seem. Unmodified GCC suffers severe bufferbloat during these events, racking up more than 1,000 frame freezes in a single five-minute session, severe enough to make interactive RTC effectively unusable over LEO links today.

To close that gap, we built PREBAS, an edge-native predictive bandwidth scaling framework driven by an ultra-lightweight 1D-CNN that recognizes the correlated RTT-creep, jitter-amplification, and loss-burst signature preceding a handoff by 200–500 ms and scales bitrate down ahead of the crash. Across 30 independent trials, the architecture cuts WebRTC frame freezes by 79% relative to GCC, lifts ITU-T P.910 QoE score by 8.36 points (43% relative) to within 2.9% of the theoretical Oracle bound, costs only 8.9% of available bandwidth, and runs inference in under 0.1 ms on commodity ARM hardware with effectively no computational penalty. At sub-millisecond latency and a 9.3 KB memory footprint, this is deployable on hardware that already exists. As LEO constellations move toward mission-critical roles, transport protocols need to anticipate degradation rather than merely recover from it. The channel emulator, training pipeline, and trained model weights will be made available in an open repository under an MIT license upon publication.

References

- [1] G. Carlucci, L. De Cicco, S. Holmer, and S. Mascolo, "Analysis and design of the Google congestion control for Web real-time communication (WebRTC)," in Proc. 7th ACM Int. Conf. Multimedia Syst. (MMSys), Klagenfurt, Austria, May 2016, pp. 13:1–13:12. doi: 10.1145/2910017.2910605.
- [2] I. Johansson, "Self-clocked rate adaptation for multimedia," IETF, RFC 8298, Dec. 2017.
- [3] S. Floyd, "Congestion control principles," IETF, RFC 2914, Sep. 2000.
- [4] G. Carlucci, L. De Cicco, S. Holmer, and S. Mascolo, "Congestion control for Web real-time communication," IEEE/ACM Trans. Netw., vol. 25, no. 5, pp. 2629–2642, Oct. 2017. doi: 10.1109/TNET.2017.2703615.
- [5] J. Zhao, O. Perrin, A. Ahangarpour, and J. Pan, "Measuring the OneWeb satellite network," in Proc. IFIP/IEEE Netw. Traffic Meas. Anal. Conf. (TMA), 2025.
- [6] M. S. B. Farooq, H. L. Vu, L. H. Andrew, and R. W. Heath, "Statistical characterization and prediction of E2E latency over LEO satellite networks," arXiv:2601.08439, 2025.
- [7] 3GPP, "NR; NR and NG-RAN overall description; Stage-2," 3GPP, TS 38.300, Rel. 17/18, 2022–2024.
- [8] A. Valentine and G. Parisi, "LeoTCP: Low-latency and high-throughput data transport for LEO satellite networks," arXiv:2508.19067, 2025.
- [9] H. Mao, R. Netravali, and M. Alizadeh, "Neural adaptive video streaming with Pensieve," in Proc. ACM SIGCOMM, Los Angeles, CA, USA, Aug. 2017, pp. 197–210. doi: 10.1145/3098822.3098843.
- [10] V. K. Adhikari et al., "Unreeling Netflix: Understanding and improving multi-CDN movie delivery," in Proc. IEEE INFOCOM, Orlando, FL, USA, Mar. 2012, pp. 1620–1628.
- [11] S. Zhang et al., "Bi-LSTM model with time distribution for bandwidth prediction in mobile networks," ETRI J., vol. 45, no. 4, pp. 612–625, Aug. 2023. doi: 10.4218/etrij.2022-0459.
- [12] M. A. Habib, P. E. Iturria-Rivera, Y. Ozcan, M. Elsayed, M. Bavand, R. Gaigalas, and M. Erol-Kantarci, "Transformer-based wireless traffic prediction and network optimization in O-RAN," in Proc. 2024 IEEE Int. Conf. Commun. Workshops (ICC Workshops), 2024, pp. 1–6.
- [13] L. Liu and Z. Xu, "Optimizing lightweight neural networks for efficient mobile edge computing," Sci. Rep., vol. 15, no. 1, art. no. 22056, Jul. 2025. doi: 10.1038/s41598-025-04652-7.
- [14] ITU-R, "Propagation data and prediction methods required for the design of Earth-space telecommunication systems," ITU-R, Rec. P.618-14, 2023.
- [15] L. De Cicco, G. Carlucci, and S. Mascolo, "Experimental investigation of the Google congestion control for real-time flows," in Proc. ACM SIGCOMM Workshop Future Human-Centric Multimedia Netw., Hong Kong, China, Aug. 2013, pp. 21–26.
- [16] Y. Zhu et al., "A comparative evaluation of TCP congestion control schemes over low-Earth-orbit (LEO) satellite networks," in Proc. ACM CoNEXT Workshop NINeS, Paris, France, Dec. 2023, pp. 1–8. doi: 10.1145/3630590.3630603.
- [17] S. Holmer et al., "A Google congestion control algorithm for real-time communication," IETF, Internet-Draft draft-ietf-rmcat-gcc-02, Oct. 2016.
- [18] A. Gottipati and L. Qiu, "Dynamic pacing for real-time satellite traffic," arXiv:2507.09798, 2025. (Preprint, under peer review at time of submission.)
- [19] J. Zhao and J. Pan, "LENS: A LEO satellite network measurement dataset," in Proc. ACM Multimedia Syst. Conf. (MMSys), 2024, pp. 278–284. doi: 10.1145/3625468.3652170.
- [20] M. Zhang, Z. Zhang, X. Liu, L. Deng, L. Zhang, Q. Gao, and H. Zhou, "Supporting real-time transmission in ultradense LEO satellite networks: A prediction-based multipath QUIC scheduler," IEEE Internet of Things J., vol. 13, no. 3, pp. 3691–3704, Feb. 2026. doi: 10.1109/JIoT.2025.3630837.
- [21] J. Zhao, J. Baude, A. Ahangarpour, V. K. Devulapalli, S. G. L. Divakarla, Z.-L. Zhang, and J. Pan, "Demystifying Starlink network performance under vehicular mobility with dynamic beam switching," arXiv:2601.13790, 2026.